

AI Engineering Intern

Applied AI, Product Engineering · Gurgaon / Hybrid

Why Bynd

Bynd is building the intelligence layer for financial services. We work with investment banks, private equity firms, asset managers, lenders, and advisory teams to transform how they extract, analyze, and act on information buried across financial documents, filings, reports, and internal workflows.

Our products sit close to real financial workflows: document intelligence, extraction, retrieval, multi-document analysis, monitoring, reporting, and workflow automation. We operate with the standards of a research team and the urgency of a product company. We care deeply about technical quality, product judgment, and building systems that get used by serious users.

The Role

As an AI Engineering Intern at Bynd, you will work closely with the engineering and product team on applied AI systems that power our platform. This is a hands-on engineering internship focused on real product work, not isolated intern projects.

You will contribute to areas such as document parsing, LLM-powered extraction, retrieval, agentic workflows, evaluation, and internal tooling. You will not be expected to own core systems independently from day one, but you will be expected to learn quickly, ship useful work, and operate with a high degree of ownership.

What you will work on

- Build and improve components in Python and TypeScript across backend services, frontend workflows, and internal tools.
- Work with LLMs and vision-language models for financial document extraction, OCR and layout interpretation, structured outputs, and workflow automation.
- Help develop prompts, schemas, evaluation sets, and validation logic that improve accuracy and reliability.
- Prototype and test RAG flows, search and retrieval pipelines, reranking logic, and multi-step AI workflows.

- Assist in turning product ideas into working features, from quick prototypes to production-ready improvements.
- Use generative AI tools thoughtfully in day-to-day work, including Claude Code, Cursor, Codex, model APIs, and internal agents.
- Work with real PDFs, filings, reports, tables, charts, and messy enterprise documents where accuracy and context matter.

What you will need

Must-haves

- Strong programming fundamentals, especially in Python. Comfort with TypeScript, React, or backend development is a strong plus.
- Genuine experience building with LLMs, not just using them. You should understand prompting, structured outputs, context windows, tool use, and common failure modes.
- Comfort using generative AI as part of your engineering workflow, while still applying your own judgment to architecture, quality, and correctness.
- Ability to move across the stack when needed. You may touch data pipelines, backend APIs, frontend flows, evaluations, and product prototypes.
- Strong attention to detail. Our users care about accuracy, citations, formatting, and reliability.
- Clear communication and the ability to work independently once direction is given.

Preferred

- Experience with RAG, embeddings, reranking, hybrid search, or evaluation frameworks.
- Experience working with OCR, document parsing, vision-language models, tables, PDFs, or multimodal AI.
- Exposure to cloud, containers, APIs, databases, CI/CD, or production engineering workflows.
- Interest in financial services workflows such as investment banking, private equity, credit, research, diligence, or reporting.
- Evidence of strong personal projects, hackathons, open-source work, research projects, or production internships.

Who you are

- You are sharp, curious, and willing to work through ambiguity.
- You like building useful products, not just experiments.

- You are comfortable with fast feedback cycles and high standards.
- You learn new tools quickly and use AI to increase your output without outsourcing your judgment.
- You care about how financial professionals actually work and why accuracy matters in context.
- You take ownership of the task in front of you and follow through.

What you will get

- Direct exposure to a fast-moving AI company building for serious financial institutions.
- Close collaboration with engineers and product folks shipping document intelligence, retrieval, extraction, and agentic workflows end-to-end.
- The chance to ship real product work that reaches users.
- A practical understanding of how modern generative AI systems are built, evaluated, and deployed.

We are looking for interns who are serious about applied AI, product engineering, and learning from real users. The role is best suited for someone who wants to build, iterate, and understand what it takes to ship reliable AI systems in production.